



# Deep Learning

## Natural Language Processing

Hugo Boulanger

M2 BDMA, CentraleSupélec, Université Paris Saclay

October 26, 2024



# Plan

- 1 Introduction
- 2 Introduction to NLP concepts



## Resources and Disclaimer

- This course is based on my experience as an NLP Researcher on what is necessary to understand the current landscape of NLP.
- I have a computer engineering background, and as such, I will not cover linguistics in depth.
- If you want to both broaden and deepen your knowledge of NLP, you can take a look Speech and Language Processing from [Jurafsky and Martin, 2024].
- This course is an evolved version of Julien Denize and Tom Dupuis's 2022 course.

# Why do we use Natural Language Processing (NLP) ?

- Languages are used for a **lot of different things**:
  - communicate: oral or written.
  - think: words shape our mind.
  - make science (maths, computer science, ...).
  - ...
- NLP is used on **various tasks**:
  - Communicate (translation, ASR, ...).
  - Make requests (search engine, recommender system, generative models).
  - Linguistics and Cognitives Sciences (Analyze language).
  - Code (Automatic code generation).
  - ...

```

26 .screen-reader-text:hover,
27 .screen-reader-text:active,
28 .screen-reader-text:focus {
29   background-color: #f1f3f4;
30   border-radius: 3px;
31   box-shadow: 0 0 2px 2px rgba(0, 0, 0, 0.6);
32   clip: auto !important;
33   color: #212529;
34   display: block;
35   font-size: 14px;
36   font-size: 0.875rem;
37   font-weight: bold;
38   height: auto;
39   left: 5px;
40   line-height: normal;
41   padding: 15px 23px 14px;
42   text-decoration: none;
43   top: 5px;
44   width: auto;
45   z-index: 100000; /* Above WP toolbar. */
46 }
47

```






## Data for Natural Language Processing

- Theoretically, the whole textual data from the web can be used.
- This represents an **enormous amount** of data.
  - 1 billion website online
  - around 60M Wikipedia pages in several languages
- Data is **increasing** at a fast rate:
  - 9000 tweets / seconds
  - 3M mails / second (60% are spam)
- Have a look yourself: <https://www.internetlivestats.com>

## Business interest for Natural Language Processing

- Search engines, with 2+billions of users for Google
- Social media with 3+billions of users (Facebook, Twitter, Instagram, TikTok, ...)
- Voice assistants with 100M+ users (Google Nest, Cortana, Alexia, Siri, ...)
- Translations (DeepL, Google translate, ...)
- Grammar check (Grammarly, Reverso, ...)

Is  grammarly  
the Right Tool for You?



# Natural Language Processing: definition and difficulty

## Natural Language Processing

**Natural language processing (NLP)** is the field of science and engineering that studies computational approaches to **understand and generate human languages**.

*"It's difficult to extract sense from strings, but they're the only communication coin we can count on. [...] Within a computer, natural language is unnatural." Alan Perlis*





## Challenges of NLP: Evolution

- Language is in **constant evolution** through time.
- New words (ex: youtuber, tiktok, ...), new grammar, new structure (ex: SMS texts).
- NLP models should **keep-up the evolution** of language which is quite difficult and require constant training.



## Challenges of NLP: Ambiguity

- Language is ambiguous, for humans and even more so for machine:
  - Meanings **depend** on the context.
  - Spelling errors.
  - ...
- Humans can ask for the correct interpretation or use **global context** (linguistic or not) to better communicate whereas machines cannot for each task.






## Challenges of NLP: Variation

- What does **variability** affect ?
  - Accent: British, indian or american
  - Spelling: color ou colour ?
  - Syntactic and Semantic, ...
- Variability comes from **several factors**:
  - **Social context**: different meaning of words depending the social class.
  - **Geography**: variation given the country.
  - **Date**: evolution of the language through time.
  - **Topic**: different meanings of words.



## Challenges of NLP: Diversity

- 7000+ languages spoken in the world
- 60% of these languages exist in written form:
  - no written data for 40% of them.
  - how to deal with few data ?
  - learning a language A can be really hard in comparison with a language B





# Plan

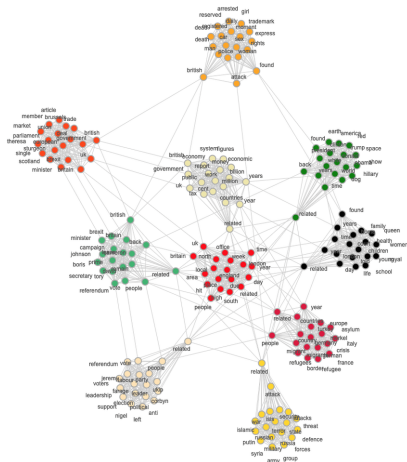
- 1 Introduction
- 2 Introduction to NLP concepts

# Corpus

## Corpus


A collection of written texts, especially the entire works of a particular author or a body of writing on a particular subject.

- In corporuses, some words come out more than others
- We can easily find **interactions** between words by:
  - calculating their co-occurrences.
  - constructing topic graphs.
  - ...



## Word distribution in a Corpus

- Word distribution of a 800 scientific articles corpus.



## Statistical Description of a Corpus

- In a large corpus, the word distributions follows the **Zipf Law**
- For  $f_W$  the frequency of the word  $w$  and  $k$  the frequency rank of the word  $w$

$$f_W(k) \propto \frac{1}{k^\theta}$$

- Most frequent words are **exponentially more frequent** than less frequent words.
- Consequences:
  - **Sparsity** among words: a lot of words have few occurrences.
  - Some very important words are less present than non-important words (such as 'a', 'the').
  - This make NLP tasks more difficult to learn.



# Token

## Token

*"A **token** is an instance of a sequence of characters in some particular document that are grouped together as a useful **semantic unit** for processing" Stanford*

- A **semantic unit** in our case can be:
  - A word or sub-word (ex: OMG)
  - A character
  - Any sequence of characters (words, multiple words, sub-words, ...)

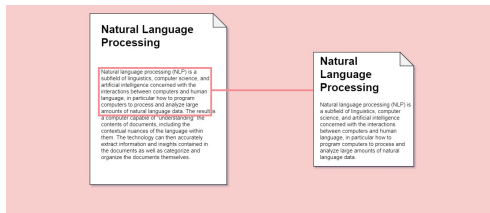


# Document

## Document

*A **document** is a sequence of semantics units.*

- The definition of document depends on the task and context and can have various length.
- Exemple of document types:
  - A tweet of 140 characters.
  - A book.
  - ...





# Corpus for NLP

## Corpus

*A **Corpus** is a collection of several documents.*

- It serves as an **input** for training a model:
  - Various size (large, medium, small), the bigger the better !
  - It can be **unbalanced**: fake news vs trustworthy news
  - It can be **messy**: typos, artefacts, ...
- It serves to various NLP tasks:
  - sentiment analysis.
  - topic extractor.
  - ...



## Tokenization

### Tokenization

***Tokenization** consists in cutting into pieces a document in **tokens**.*

- Some characters can be **thrown away** sometimes like punctuation or whitespaces.
- What are the tokens to use ?
  - Some words might not be important: "the".
  - Some words only appear once.
  - What about apostrophes: "O'neill" becomes "o" + "neill" or "o'" + "neill" or ...
- Tokenization is **language specific**.
- A class of token is sometimes called a **type**.
- The different types kept form the **vocabulary** or **index** of the corpus.



# Language Modeling

## Language Modeling

Language Modeling is the process of assigning probabilities to the next word(s) (or, in general the next token(s)).

## Language Model

A language model, is a model assigning probabilities to the possible next word.

Given tokens, how do we build a language model?



## Language Modeling: the maths

### N-gram

An N-gram is a tuple of N words (ordered).

- For  $w_n$  the  $n^{th}$  word in the sentence:  
$$P(w_{1:n}) = P(w_1)P(w_2|w_1)...P(w_n|w_{1:n-1})$$
$$P(w_{1:n}) = \prod_{i=1}^n P(w_i|w_{1:i-1})$$
- We can approximate using the Markov assumption (N-gram):  
$$P(w_n|w_{1:n-1}) = P(w_n|w_{n-N+1:n-1})$$
- Now :  $P(w_{1:n}) = \prod_{i=1}^n P(w_i|w_{n-N+1:i-1})$



## N-gram Models

N-gram models build their probabilities by counting the occurrences of the N-grams:

$$P(w_n | w_{n-N+1:n-1}) = \frac{C(w_{n-N+1:n})}{\sum_w C(w_{n-N+1:n-1}w)}$$

The denominator is equal to the count of the N-1-gram

$$P(w_n | w_{n-N+1:n-1}) = \frac{C(w_{n-N+1:n})}{C(w_{n-N+1:n-1})}$$



# Word Embeddings

## Word Embeddings

**Word vectors** are vectors representing words. They can also be called **features or representations of words**.

- Word Embeddings allow for a dense representation of words instead of discrete values from tokenization.
- Tokens are at word-level (can be sub-word level, as we'll see later)



## Prediction-based model

- **Learn dense vectors** to represent words through an **embedding matrix** later used to **convert** tokens into vectors.
- **Distributional Hypothesis**: use context to build the vectors.
- Parametrize words as dense vectors.
- Use parametrization to **predict the context** and learn the representation.





## Self-Supervised Word2vec

- Word2vec **maps** each word to a vector of a certain dimension.
- It can be made from **two models** [Mikolov et al., 2013] detailed after:
  - Skip-gram.
  - Continuous Bag of Words (CBoW).
- The training seeks to **predict words in the context**, defined as surrounding words, from a word.
- Therefore Word2vec **does not require labels** and uses its own data for supervision, also called self-supervised learning.

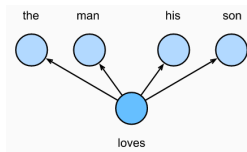
## Skip-Gram

- Take the sentence: "The", "man", "loves", "his", "son".
- The center word is defined as "loves" and the context window is 2 (all other words for this specific sentence).
- Skip-Gram seeks to predict the conditional probability to **generate** the context:

$$P("the", "man", "his", "son" | "loves")(1)$$

- It assumes that context words are samples independently, therefore:

$$(1) = P("the" | "loves) \cdot P("man" | "loves) \cdot P("his" | "loves) \cdot P("son" | "loves)$$



## Skip-Gram

- Each word  $w_i$  is associated to **two learned** representations:
  - $v_i \in \mathbb{R}^d$  **center** word vector
  - $u_i \in \mathbb{R}^d$  **context** word vector
- The probability to generate the context word  $w_o$  **conditionally** to the word  $w_c$  from the vocabulary index set  $\mathcal{V} = \{0, 1, \dots, |\mathcal{V}| - 1\}$ , is defined as the following **softmax** operation:

$$P(w_o|w_c) = \frac{\exp(\mathbf{u}_o^T \mathbf{v}_c)}{\sum_{i \in \mathcal{V}} \exp(\mathbf{u}_i^T \mathbf{v}_c)}$$

- All words used in denominator are considered as **negatives** that the true context word, **positive**, should be contrasted with.
- Contrastive learning in images took inspiration from this !

## Skip-Gram

- The probability to generate the context word  $w_o$  **conditionally** to the word  $w_c$  is:

$$P(w_o|w_c) = \frac{\exp(\mathbf{u}_o^T \mathbf{v}_c)}{\sum_{i \in \mathcal{V}} \exp(\mathbf{u}_i^T \mathbf{v}_c)}$$

- For a context window of size  $m$  the **likelihood function** is the probability to generate all context words from a text sequence of length  $T$  as follows:

$$\prod_{t=1}^T \prod_{-m \leq j \leq m, j \neq 0} P(w^{t+j}|w^t)$$

- For training we seek to **maximize the log-likelihood function** of the skip-gram model which is equivalent to minimizing:

$$-\sum_{t=1}^T \sum_{-m \leq j \leq m, j \neq 0} \log P(w^{t+j}|w^t)$$

## Skip-Gram: training

- For training we seek to **maximize the log-likelihood function** of the skip-gram model which is equivalent to minimizing:

$$-\sum_{t=1}^T \sum_{-m \leq j \leq m, j \neq 0} \log P(w^{t+j} | w^t)$$

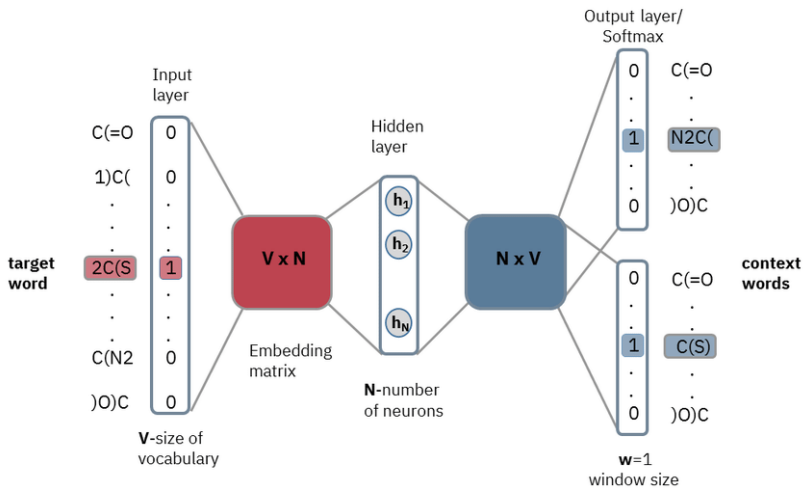
- We need to compute each log conditional probability for any pair  $w_c$  and  $w_o$ :

$$\log P(w_o | w_c) = \mathbf{u}_o^T \mathbf{v}_c - \log \left( \sum_{i \in \mathcal{V}} \exp(\mathbf{u}_i^T \mathbf{v}_c) \right)$$

- To train using **gradient-descent**, we compute the **partial derivatives**:

$$\frac{\partial \log P(w_o | w_c)}{\partial v_c}, \frac{\partial \log P(w_o | w_c)}{\partial u_o}, \forall j, \frac{\partial \log P(w_o | w_c)}{\partial u_j}$$

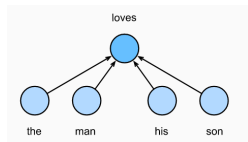
## Skip-Gram: illustration



## Continuous Bag of Words (CBOW)

- Similar to Skip-Gram model.
- Continuous Bag of Words (CBOW) seeks to predict the conditional probability to **generate** the center word given its context is:

$$P(\text{"loves"} | \text{"the"}, \text{"man"}, \text{"his"}, \text{"son"})$$



- The conditional probability to generate the word  $w_c$  given its context  $w_{o_1}, \dots, w_{o_{2m}}$  is:

$$P(w_c | w_{o_1}, \dots, w_{o_{2m}}) = \frac{\exp(\frac{1}{2m} \mathbf{u}_c^T (\mathbf{v}_{o_1} + \dots + \mathbf{v}_{o_{2m}}))}{\sum_{i \in \mathcal{V}} \exp(\frac{1}{2m} \mathbf{u}_i^T (\mathbf{v}_{o_1} + \dots + \mathbf{v}_{o_{2m}}))}$$



## Continuous Bag of Words (CBOW)

- The conditional probability to generate the word  $w_c$  given its context  $w_{o_1}, \dots, w_{o_{2m}}$ :

$$P(w_c | w_{o_1}, \dots, w_{o_{2m}}) = \frac{\exp(\frac{1}{2m} \mathbf{u}_c^T (\mathbf{v}_{o_1} + \dots + \mathbf{v}_{o_{2m}}))}{\sum_{i \in \mathcal{V}} \exp(\frac{1}{2m} \mathbf{u}_i^T (\mathbf{v}_{o_1} + \dots + \mathbf{v}_{o_{2m}}))}$$

- Let's define  $v_o = \frac{1}{2m} (v_{o_1} + \dots + v_{o_{2m}})$  then:

$$P(w_c | w_{o_1}, \dots, w_{o_{2m}}) = \frac{\exp(\mathbf{u}_c^T \mathbf{v}_o)}{\sum_{i \in \mathcal{V}} \exp(\mathbf{u}_i^T \mathbf{v}_o)}$$





## Continuous Bag of Words (CBOW) training

- For a context window of size  $m$  the likelihood function is the probability to **generate** the center word given its context from a text sequence of length  $T$ , and defined as follows:

$$\prod_{t=1}^T P(w^t | w^{t-m}, \dots, w^{t-1}, w^{t+1}, \dots, w^{t+m})$$

- For training we seek to maximize the log-likelihood function of the CBOW model which is equivalent to minimizing:

$$- \sum_{t=1}^T \log P(w^t | w^{t-m}, \dots, w^{t-1}, w^{t+1}, \dots, w^{t+m})$$

with

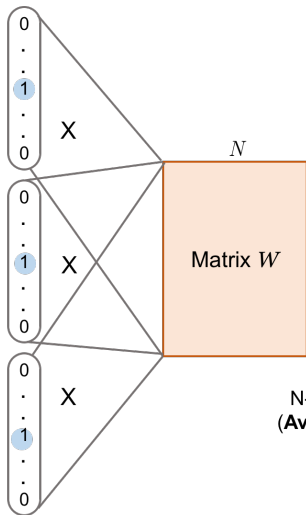
$$\log(P_{w_c} | w_{o_1}, \dots, w_{o_{2m}}) = \mathbf{u}_c^T \mathbf{v}_o - \log \left( \sum_{i \in \mathcal{V}} \exp(\mathbf{u}_i^T \mathbf{v}_o) \right)$$

- To train using **gradient-descent** we compute the partial derivatives.

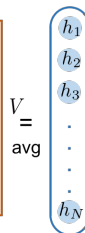


## CBoW: illustration


Input



Hidden



N-dimension vector  
(**Average** of vectors of  
all input words)

Output  
softmax



## Training Word2vec

- Word2vec output is the matrix embedding from **either** Skip-gram or CBoW.
- In practice CBoW is quicker to train but Skip-gram has better results.
- Improvements:
  - **Negative sampling**: draw negatives from a predefined distribution.
  - **Hierarchical softmax**: uses binary tree.
  - **Global Statistics**: Additional use of global statistics led to GloVe[Pennington et al., 2014]
- Problems:
  - **Fixed Vocabulary**: unknown tokens are not treated.
  - **Fixed Context**: static embeddings don't react to context.

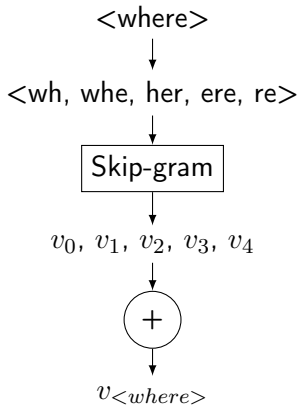





## Solving Fixed Vocabulary

### Words

Words are made of sub-words.

- Words are sequences of **n-grams of characters**.
- You can learn a model on this sequence instead.
- You obtain word representation by summing the sub-word representations.
- This technique is called **Fast-Text** [Bojanowski et al., 2016]



-  Bojanowski, P., Grave, E., Joulin, A., and Mikolov, T. (2016).  
Enriching word vectors with subword information.  
*arXiv preprint arXiv:1607.04606*.
-  Jurafsky, D. and Martin, J. H. (2024).  
*Speech and Language Processing: An Introduction to Natural Language Processing, Computational Linguistics, and Speech Recognition with Language Models*.  
Stanford University, 3rd edition.  
Online manuscript released August 20, 2024.
-  Mikolov, T., Sutskever, I., Chen, K., Corrado, G. S., and Dean, J. (2013).  
Distributed representations of words and phrases and their compositionality.  
*In 27th Annual Conference on Neural Information Processing Systems*, pages 3111–3119.



Pennington, J., Socher, R., and Manning, C. D. (2014).  
Glove: Global vectors for word representation.  
In *EMNLP*, pages 1532–1543.